

# Lecture 7: Dynamic Relocation and Segmentation

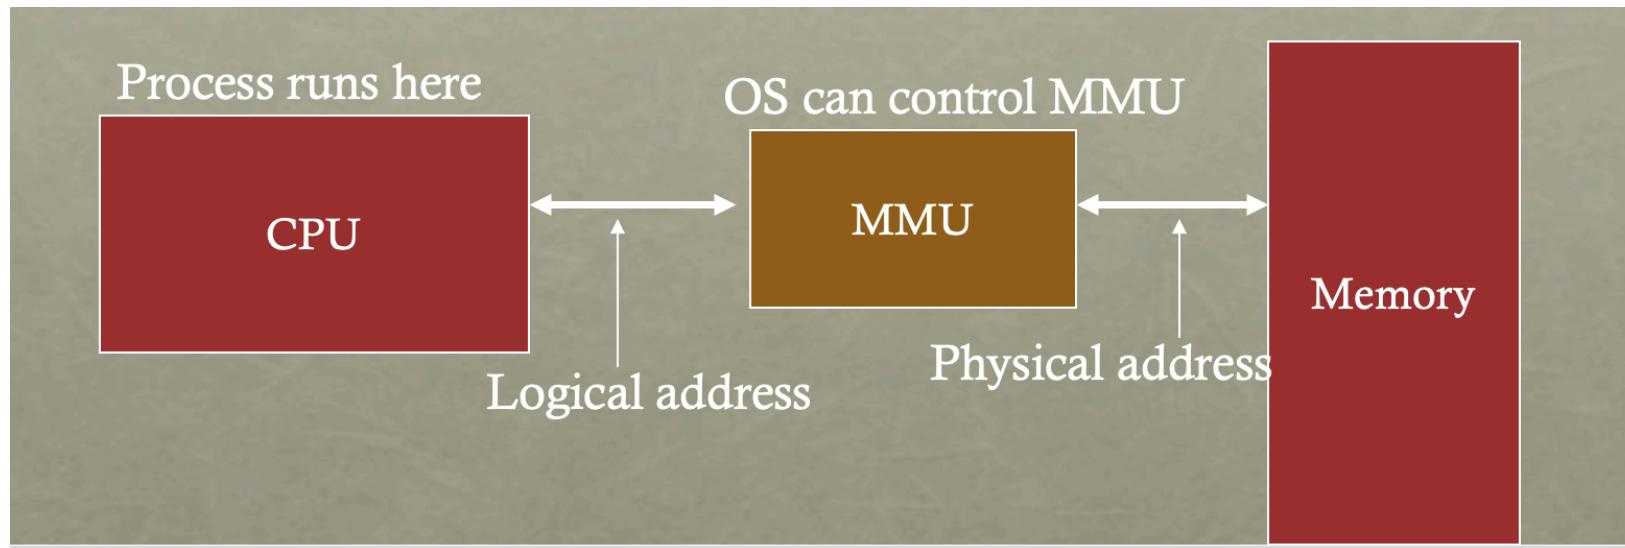
Operating Systems

**Content taken from:** <https://www.cse.iitb.ac.in/~mythili/os/>  
<https://pages.cs.wisc.edu/~remzi/OSTEP/>

# Last Class

- **Virtual Address Space:** Every process assumes it has access to a large space of memory from address 0 to a MAX
- CPU issues loads and stores to virtual addresses
- Address translation is required to go from virtual addresses (VA) to physical addresses (PA)
- **Memory Management Unit (MMU):** Hardware which performs the address translation from VA to PA

# Address Translation using MMU



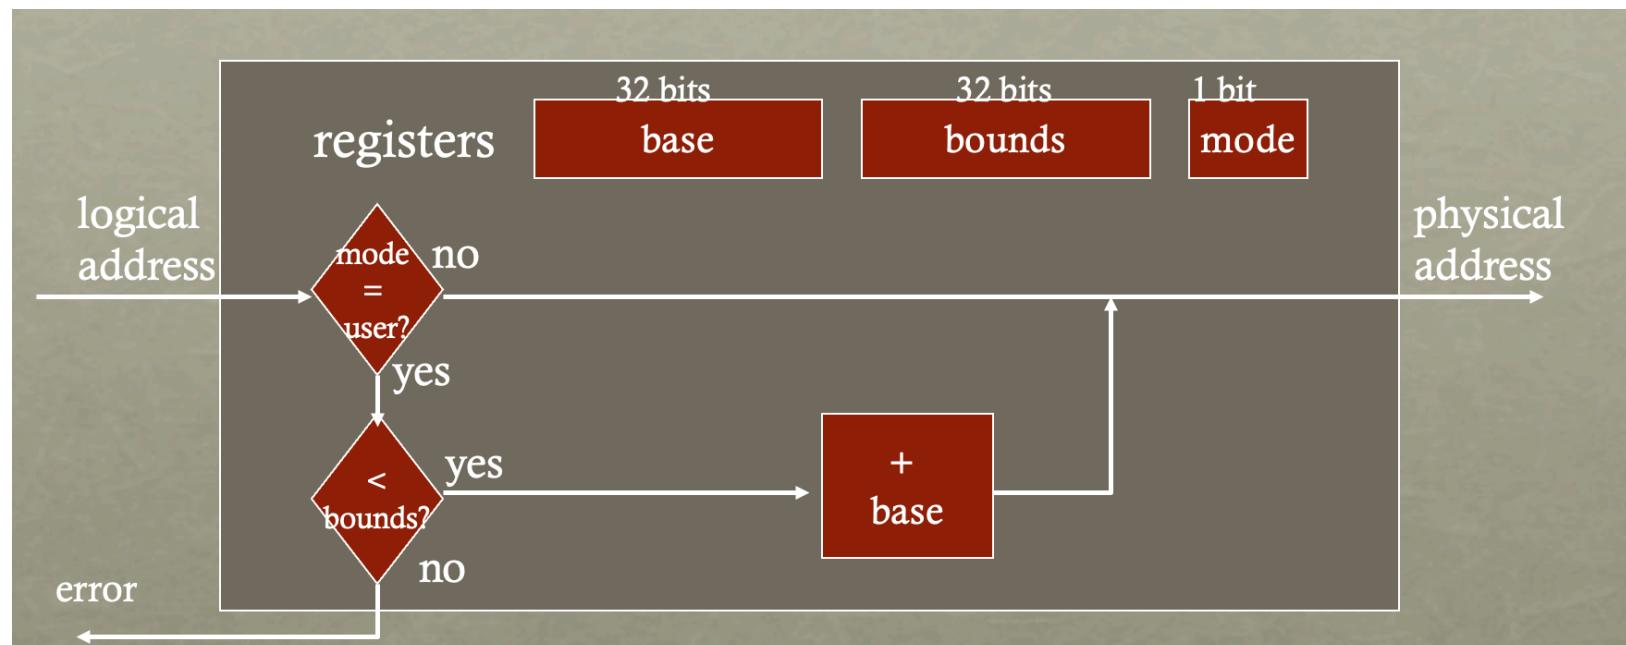
# Dynamic Relocation with Base and Bounds

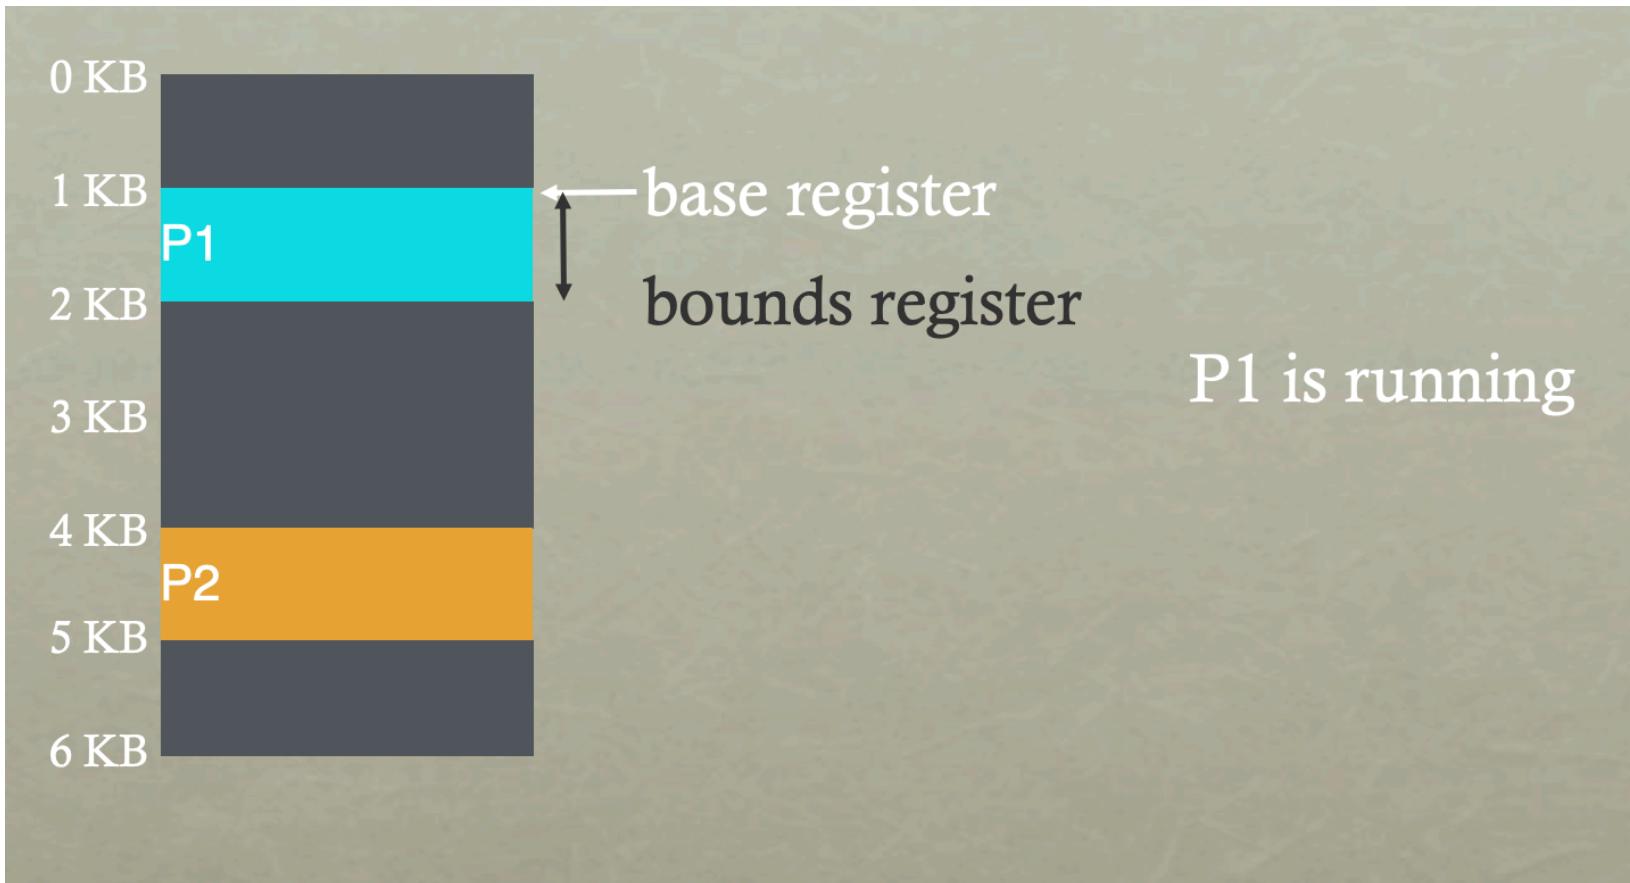
- **Base register:** Contains the offset to be added to VA to get the PA
  - smallest physical address (or starting location)
- **Bounds register:** Ensures that all virtual addresses generated by the process are legal and within the “bounds” of the process.
  - size of this process’s virtual address space

# Implementation of Base + Bounds

Translation on every memory access of user process

- MMU compares logical address to bounds register
  - if logical address is greater, then generate error
- MMU adds base register to logical address to form physical address









Can P1 hurt P2?

Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	

P2: load 1000, R1  
 $\Rightarrow 4096 + 1000 =$   
 5096 and not 5196

P1: load 1000, R1  
 $\Rightarrow 1024 + 1000 =$   
 2024. It should be  
 load 1000, R1 and  
 not load 100, R1.



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	interrupt OS! 3072 > 1024

P2: load 1000, R1  
 $\Rightarrow 4096 + 1000 =$   
 5096 and not 5196

P1: load 1000, R1  
 $\Rightarrow 1024 + 1000 =$   
 2024. It should be  
 load 1000, R1 and  
 not load 100, R1.



Can P1 hurt P2?

	Virtual	Physical
0 KB		
1 KB	P1: load 100, R1	load 1124, R1
2 KB	P2: load 100, R1	load 4196, R1
3 KB	P2: load 1000, R1	load 5196, R1
4 KB	P1: load 100, R1	load 2024, R1
5 KB	P1: store 3072, R1	interrupt OS!

P2: load 1000, R1  
 $\Rightarrow 4096 + 1000 =$   
 5096 and not 5196

P1: load 1000, R1  
 $\Rightarrow 1024 + 1000 =$   
 2024. It should be  
 load 1000, R1 and  
 not load 100, R1.

# Managing Processes with Base and Bounds

## **Context-switch**

- Store base and bounds registers in PCB (Process Control Block)
- Steps
  - Change to privileged mode
  - Save base and bounds registers of old process
  - Load base and bounds registers of new process
  - Change to user mode and jump to new process

**What if don't change base and bounds registers when switching?**

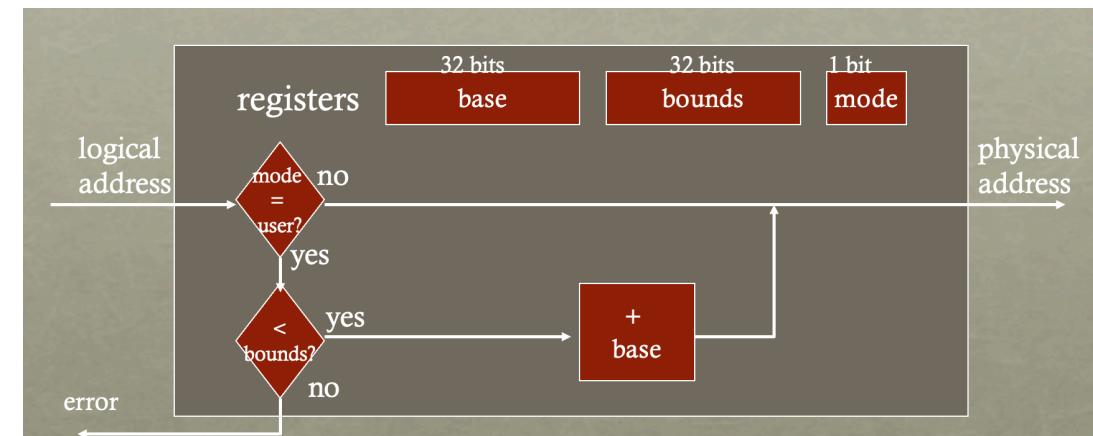
## **Protection requirement**

- User process cannot change base and bounds registers
- User process cannot change to privileged mode

# Base and Bounds Advantages

## Advantages

- Provides protection (both read and write) across address spaces
- Supports dynamic relocation
  - Can place process at different locations initially and also move address spaces
- Simple, inexpensive implementation
  - Few registers, little logic in MMU
- Fast
  - Add and compare in parallel



# Base and Bounds Disadvantages

## Disadvantages

- Each process must be allocated contiguously in physical memory
  - Must allocate memory that may not be used by process
- No partial sharing: Cannot share limited parts of address space

<b>Hardware Requirements</b>	<b>Notes</b>
Privileged mode	<i>Needed to prevent user-mode processes from executing privileged operations</i>
Base/bounds registers	<i>Need pair of registers per CPU to support address translation and bounds checks</i>
Ability to translate virtual addresses and check if within bounds	<i>Circuitry to do translations and check limits; in this case, quite simple</i>
Privileged instruction(s) to update base/bounds	<i>OS must be able to set these values before letting a user program run</i>
Privileged instruction(s) to register exception handlers	<i>OS must be able to tell hardware what code to run if exception occurs</i>
Ability to raise exceptions	<i>When processes try to access privileged instructions or out-of-bounds memory</i>

Figure 15.3: Dynamic Relocation: Hardware Requirements

OS Requirements	Notes
Memory management	<i>Need to allocate memory for new processes; Reclaim memory from terminated processes; Generally manage memory via free list</i>
Base/bounds management	<i>Must set base/bounds properly upon context switch</i>
Exception handling	<i>Code to run when exceptions arise; likely action is to terminate offending process</i>

Figure 15.4: Dynamic Relocation: Operating System Responsibilities

OS @ boot (kernel mode)	Hardware	(No Program Yet)
initialize trap table	remember addresses of... system call handler timer handler illegal mem-access handler illegal instruction handler	
start interrupt timer		start timer; interrupt after X ms
initialize process table		
initialize free list		

Figure 15.5: Limited Direct Execution (Dynamic Relocation) @ Boot

OS @ run (kernel mode)	Hardware	Program (user mode)
To start process A: allocate entry in process table alloc memory for process set base/bound registers <b>return-from-trap</b> (into A)	restore registers of A move to <b>user mode</b> jump to A's (initial) PC	<b>Process A runs</b> Fetch instruction
	translate virtual address perform fetch	Execute instruction
	if explicit load/store: ensure address is legal translate virtual address perform load/store	
	<b>Timer interrupt</b> move to <b>kernel mode</b> jump to handler	(A runs...)
<b>Handle timer</b> decide: stop A, run B call <code>switch()</code> routine save regs(A) to proc-struct(A) (including base/bounds) restore regs(B) from proc-struct(B) (including base/bounds) <b>return-from-trap</b> (into B)	restore registers of B move to <b>user mode</b> jump to B's PC	<b>Process B runs</b> Execute bad load
	Load is out-of-bounds; move to <b>kernel mode</b> jump to trap handler	
<b>Handle the trap</b> decide to kill process B deallocate B's memory free B's entry in process table		

Figure 15.6: Limited Direct Execution (Dynamic Relocation) @ Runtime

# Segmentation

Divide address space into logical segments

- Each segment corresponds to logical entity in address space
- A segment is just a contiguous portion of the address space of a particular length
- Code, stack, heap

Each segment can independently:

- be placed separately in physical memory
- grow and shrink
- be protected (separate read/write/execute protection bits)

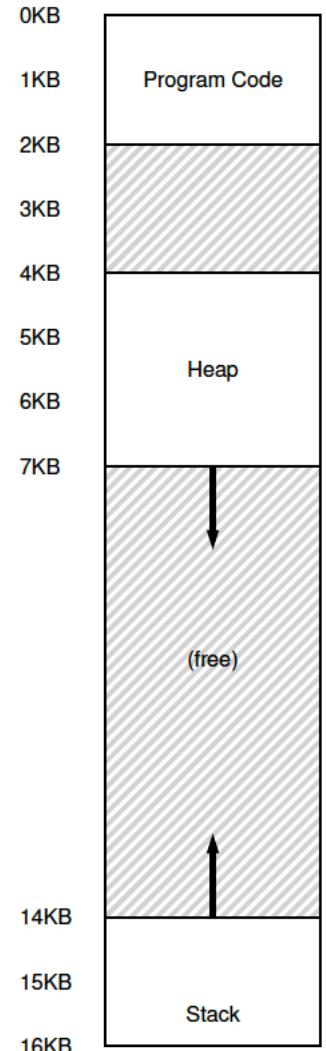
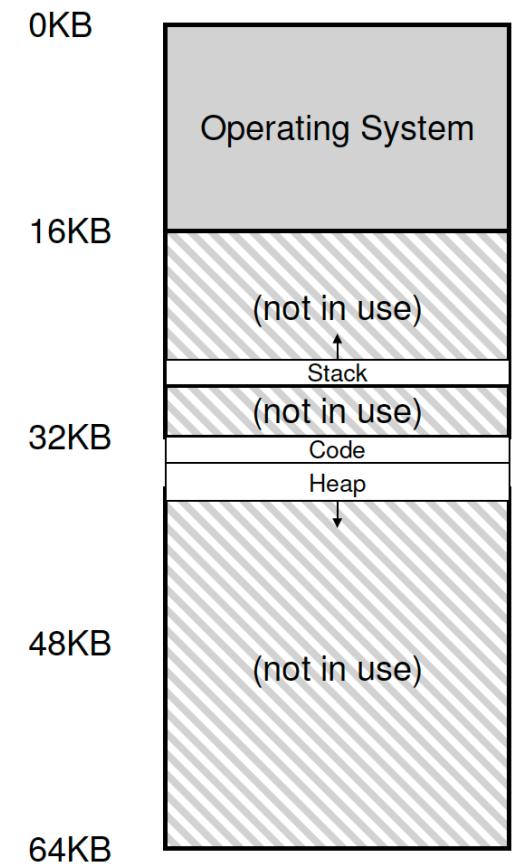


Figure 16.1: An Address Space (Again)

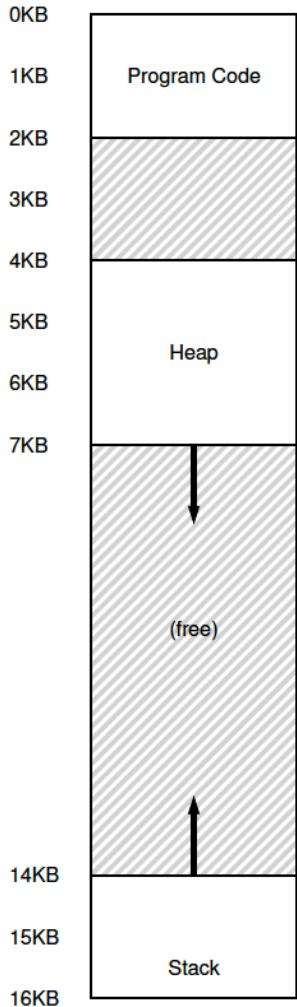
# MMU Support for Segmentation

- MMU keeps separate base and bound register values for each segment of the process

Segment	Base	Size
Code	32K	2K
Heap	34K	3K
Stack	28K	2K



# Example Translations for Segmentation



Segment	Base	Size
Code	32K	2K
Heap	34K	3K
Stack	28K	2K

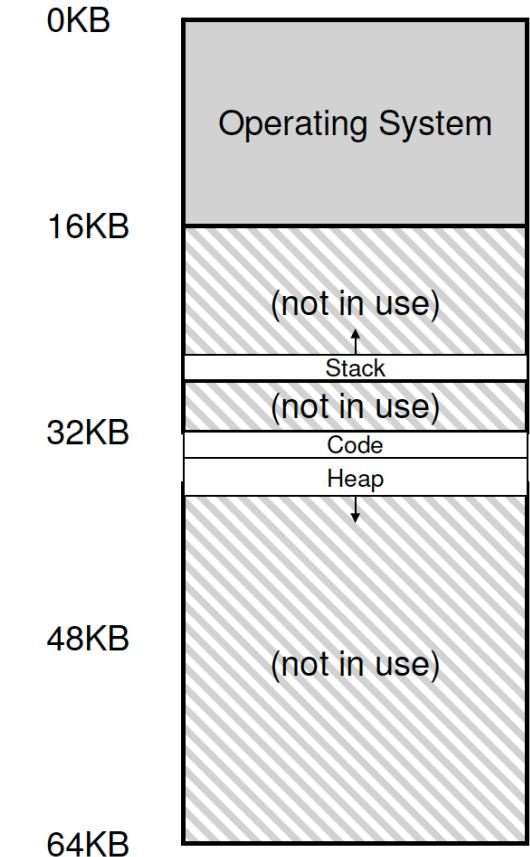
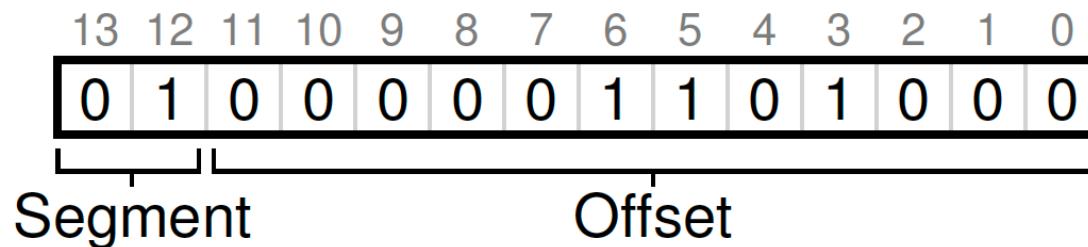


Figure 16.1: An Address Space (Again)

# Which Segment Are We Referring To?

- How does process designate a particular segment?
  - Use part of logical address
    - Top bits of logical address select segment
    - Low bits of logical address select offset within segment



# What about the stack?

- MMU keeps track of whether the segment grows in the positive or negative direction

Segment	Base	Size (max 4K)	Grows Positive?
Code <sub>00</sub>	32K	2K	1
Heap <sub>01</sub>	34K	3K	1
Stack <sub>11</sub>	28K	2K	0

Figure 16.4: Segment Registers (With Negative-Growth Support)

- Address translation for Stack is also slightly different.
- Example: 15 KB virtual address should map to ?

# Support for Sharing

- Sometimes it is useful to share certain memory segments between address spaces.
  - Code sharing is common
- Hardware also stores **protection bits** per segment

Segment	Base	Size (max 4K)	Grows Positive?	Protection
Code <sub>00</sub>	32K	2K	1	Read-Execute
Heap <sub>01</sub>	34K	3K	1	Read-Write
Stack <sub>11</sub>	28K	2K	0	Read-Write

Figure 16.5: Segment Register Values (with Protection)

# OS Support for Segmentation

- What should the OS do on a context switch?
- What happens when a segment grows or shrinks?
- How to allocate new segments or grow existing ones?
  - External Fragmentation
  - Solution: Use a free-list management algorithm
    - Best-Fit
    - Worst-Fit
    - First-fit

# Issue with Segmentation: Fragmentation

- Different-sized segments are allocated memory.
- Free memory (hole) is too small and scattered
  - Segment to be allocated may be larger than the free memory hole but smaller than the total amount of free memory available