

# Lecture 10: Swapping

Operating Systems

**Content taken from:** <https://pages.cs.wisc.edu/~remzi/OSTEP/>

# Last few classes

- MMU for translating a virtual address to physical address
- Dynamic relocation with base + bound registers
- Segmentation
- Paging
- Multi-Level Paging
- TLB (cache for virtual-to-physical address translations)

# Assumption till now

- Every virtual address space of every running process fits into physical memory.
  - That is, all the pages reside in physical memory.
- Can we relax this assumption?
  - Yes, we can move some of the pages (allocated to processes) to disk
- Why do we need to relax this assumption?
  - To support the illusion of a large virtual memory for multiple concurrently-running processes.

# Swap Space

- Reserved space on disk for moving pages back and forth
- OS can read from and write to the swap space, in page-sized units.
  - OS needs to remember the **disk address** of a given page.




Figure 21.1: Physical Memory and Swap Space

# Present Bit

- Till now, we have looked at
  - Valid bit
  - Protected bit
- Latest addition: Present bit in each PTE
  - If set to one, then the page is present in physical memory
  - If set to zero, then the page is not present in physical memory but stored on disk

# Page Fault

- Trying to access a page that is not present in physical memory
- Upon a page fault, OS is invoked to service the page fault using a **page-fault handler**
- OS finds the disk address for the page in it's PTE.
- OS issues the request to disk to fetch the page into memory
- Once the page is fetched into memory
  - OS updates the PTE to mark the page as present
  - OS updates the PFN of the PTE with the new in-memory location

# What if memory is full?

- Replace an existing page with the new pages which need to be brought in
- Page-replacement policy dictates which page to replace?
  - Will talk about it next.

```
1  PFN = FindFreePhysicalPage()
2  if (PFN == -1)          // no free page found
3      PFN = EvictPage()  // run replacement algorithm
4  DiskRead(PTE.DiskAddr, PFN) // sleep (waiting for I/O)
5  PTE.present = True     // update page table with present
6  PTE.PFN      = PFN    // bit and translation (PFN)
7  RetryInstruction()    // retry instruction
```

Figure 21.3: Page-Fault Control Flow Algorithm (Software)

```

1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True) // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10     else // TLB Miss
11         PTEAddr = PTBR + (VPN * sizeof(PTE))
12         PTE = AccessMemory(PTEAddr)
13         if (PTE.Valid == False)
14             RaiseException(SEGMENTATION_FAULT)
15         else
16             if (CanAccess(PTE.ProtectBits) == False)
17                 RaiseException(PROTECTION_FAULT)
18             else if (PTE.Present == True)
19                 // assuming hardware-managed TLB
20                 TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
21                 RetryInstruction()
22             else if (PTE.Present == False)
23                 RaiseException(PAGE_FAULT)

```

Figure 21.2: Page-Fault Control Flow Algorithm (Hardware)

# Optimal Replacement Policy

- Main memory can be viewed as a cache for virtual memory pages in the system
- Optimal replacement policy should
  - Minimize the number of cache misses
    - Minimize the number of times we have to fetch a page from disk
  - Maximize the number of cache hits
    - Maximize the number of times a page that is accessed is found in memory
- **Optimal policy:** Replace the page that will be accessed **furthest in the future**

# Optimal: Example

- Compute the cache hit rate?

- $\# \text{hits} / (\# \text{hits} + \# \text{misses})$
  - $= 6 / (6 + 5)$
  - Page 3 of [Textbook chapter URL](#)

Access	Hit/Miss?	Evict	Resulting Cache State
0			
1			
2			
0			
1			
3			
0			
3			
1			
2			
1			

# FIFO replacement policy

- Pages are stored in a queue
- Page on the tail of the queue (First-in) is evicted.

# FIFO: Example

- Compute the cache hit rate?
  - 4 hits, 7 misses
  - Hit rate =  $(4/4+7)$
  - Page 5 of [Textbook chapter URL](#)

Access	Hit/Miss?	Evict	Resulting Cache State
0			
1			
2			
0			
1			
3			
0			
3			
1			
2			
1			

# Using History: LRU

- If a process has accessed a page in the near past, it is likely to access it again in the near future
- Least-Recently-Used (LRU) policy replaces the least-recently-used page.
  - The more recently a page has been accessed, perhaps the more likely it will be accessed again
- **Least-Frequently-Used (LFU)** policy replaces the least-frequently-used page
  - If a page has been accessed many times, perhaps it should not be replaced as it clearly has some value.

# LRU: Example

- Compute the cache hit rate?
  - 6 hits, 5 misses
  - Hit rate =  $(6/6+5)$
  - Page 7 of [Textbook chapter URL](#)

Access	Hit/Miss?	Evict	Resulting Cache State
0			
1			
2			
0			
1			
3			
0			
3			
1			
2			
1			

# Implementing LRU

- **Option 1:** Maintain a list of pages where the most recently used page is moved to the front of the list on every page access.
- **Option 2:** Update the time field associated with the page just accessed. For replacement, find the page with the oldest timestamp.
- Both options are costly

# Approximating LRU

- Store a **use** bit (or **reference** bit) in each PTE
- Whenever a page is referenced (read/written), the **use** bit is set to 1
- How to employ the **use** bit to approximate LRU? [**Clock algorithm**]
  - Imagine all the pages to be arranged in a circular list
  - OS scans through this list and replaces the first page whose use bit is set to 0.
  - Whenever the OS encounters a page whose use bit is set to 1, that use bit is reset to 0.
  - The OS then proceeds until it finds a page whose use bit is set to 0.

# Considering Dirty Pages

- If a page has been modified in memory and not yet written to disk, it is marked as dirty using the **dirty** bit.
- The page replacement policy does not evict a dirty page as before eviction, the dirty page necessarily needs to be written to disk which is expensive.
- Clock algorithm can be modified to first consider pages which has both use bit and dirty bit set to 0.

# Other policies

- **Demand paging:** OS brings a page into memory when it is accessed i.e. “on demand”
- **Thrashing:**
  - Memory demands of the set of running processes simply exceeds the available physical memory
  - System will be constantly paging (out and in)