

# Lecture 14: Locks and Condition Variables

Operating Systems

**Content taken from:** <https://pages.cs.wisc.edu/~remzi/OSTEP/>  
<https://www.cse.iitb.ac.in/~mythili/os/>

# Last Class

- Process vs threads
  - Parent process spawns a child process using `fork()`.
    - No sharing of memory between the parent and child
  - Parent process spawns a new thread using `pthread_create()`
    - Parent and new thread share the same address space (except the stack)
    - Threads are like separate processes except they share the same address space
- Why threads? Parallelism and effective usage of CPU
- OS schedules threads similar to how it schedules processes

# Last Class

- Threads of the same process share data
- **Race Conditions:**
  - Two threads trying to access a **critical section** (updating a shared variable)
  - Leads to indeterminate/unexpected results
- **Mutual Exclusion:**
  - Ensure that only one thread at a time executes the critical section
  - Requires that the critical section executes **atomically** (cannot be interrupted by any interrupts etc.)
  - How? Using Locks

# Locks

- For each critical section, use a lock
- Each thread need to get the lock before executing the critical section
- At a time, only one thread can have the lock
- If the lock is held by another thread, then all other threads wait until the lock is released
- Modern architectures provide hardware atomic instructions for implementing locks

# Test-and-Set Instruction

- Update a variable and return old value, all in one hardware instruction

```
1  int TestAndSet(int *old_ptr, int new) {  
2      int old = *old_ptr; // fetch old value at old_ptr  
3      *old_ptr = new;    // store 'new' into old_ptr  
4      return old;       // return the old value  
5  }
```

# Spin lock using test-and-set

- If `TestAndSet(flag,1)` returns 1, it means the lock is held by someone else, so spin until the flag changes to 0

```
1  typedef struct __lock_t {  
2      int flag;  
3  } lock_t;  
4  
5  void init(lock_t *lock) {  
6      // 0: lock is available, 1: lock is held  
7      lock->flag = 0;  
8  }  
9  
10 void lock(lock_t *lock) {  
11     while (TestAndSet(&lock->flag, 1) == 1)  
12         ; // spin-wait (do nothing)  
13 }  
14  
15 void unlock(lock_t *lock) {  
16     lock->flag = 0;  
17 }
```

Figure 28.3: A Simple Spin Lock Using Test-and-set

# Evaluating Spin Locks

- Correctness?
- Fairness?
- Performance?
- Think about the above properties when there are  $N$  threads being scheduled in a round-robin fashion?

# How to avoid wasteful spinning?

- Can we avoid wasteful spinning while the thread is waiting for the lock to be released?
- Use **yield()**
- Yield() moves the waiting thread from running to ready state

```
1  void init() {  
2      flag = 0;  
3  }  
4  
5  void lock() {  
6      while (TestAndSet(&flag, 1) == 1)  
7          yield(); // give up the CPU  
8  }  
9  
10 void unlock() {  
11     flag = 0;  
12 }
```

# **Yield() is helpful but...**

- **Performance:** While better than spinning, yield() does have overhead due to the associated cost of context switch
- **Starvation/Fairness:** A thread may get caught in an endless yield loop while other threads repeatedly enter and exit the critical section
- **What is the problem in the approaches we discussed?**
  - Scheduler may make a bad choice: it may schedule a thread which yields the CPU immediately or spins for the entire time slice.
  - There is potential for waste and no prevention for starvation

# Using Sleep and Queues

- Use sleep instead of spinning or yielding
  - sleep moves the process from running to blocked
  - yield moves the process from running to ready
- We will have a queue to keep track of which threads are waiting to acquire the lock
- How to build a lock using sleep?
  - If the lock is held by another thread, add the calling thread to the queue and put it to sleep
  - If the lock is available, wake up the first sleeping thread from the queue

```

1  typedef struct __lock_t {
2      int flag;
3      int guard;
4      queue_t *q;
5  } lock_t;
6
7  void lock_init(lock_t *m) {
8      m->flag = 0;
9      m->guard = 0;
10     queue_init(m->q);
11 }
12
13 void lock(lock_t *m) {
14     while (TestAndSet(&m->guard, 1) == 1)
15         ; //acquire guard lock by spinning
16     if (m->flag == 0) {
17         m->flag = 1; // lock is acquired
18         m->guard = 0;
19     } else {
20         queue_add(m->q, gettid());
21         m->guard = 0;
22         park();
23     }
24 }
25
26 void unlock(lock_t *m) {
27     while (TestAndSet(&m->guard, 1) == 1)
28         ; //acquire guard lock by spinning
29     if (queue_empty(m->q))
30         m->flag = 0; // let go of lock; no one wants it
31     else
32         unpark(queue_remove(m->q)); // hold lock
33                                     // (for next thread!)
34     m->guard = 0;
35 }

```

Figure 28.9: Lock With Queues, Test-and-set, Yield, And Wakeup

# Condition Variables: Waiting and Signaling

- Locks allow one type of synchronization between threads – mutual exclusion
- Another common requirement in multi-threaded applications – waiting and signaling
  - E.g., Thread T1 wants to continue only after T2 has finished some task
- Can accomplish this by busy-waiting on some variable, but inefficient
- Need a new synchronization primitive: condition variables

```
1  volatile int done = 0;
2
3  void *child(void *arg) {
4      printf("child\n");
5      done = 1;
6      return NULL;
7  }
8
9  int main(int argc, char *argv[]) {
10     printf("parent: begin\n");
11     pthread_t c;
12     Pthread_create(&c, NULL, child, NULL); // create child
13     while (done == 0)
14         ; // spin
15     printf("parent: end\n");
16     return 0;
17 }
```

Figure 30.2: Parent Waiting For Child: Spin-based Approach

# Condition Variables

- A condition variable (CV) is a queue that a thread can put itself into when waiting on some condition
- Another thread that makes the condition true can signal the CV to wake up a waiting thread

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);  
pthread_cond_signal(pthread_cond_t *c);
```

# Example

```
1  int done = 0;
2  pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3  pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5  void thr_exit() {
6      Pthread_mutex_lock(&m);
7      done = 1;
8      Pthread_cond_signal(&c);
9      Pthread_mutex_unlock(&m);
10 }
11
12 void *child(void *arg) {
13     printf("child\n");
14     thr_exit();
15     return NULL;
16 }
17
18 void thr_join() {
19     Pthread_mutex_lock(&m);
20     while (done == 0)
21         Pthread_cond_wait(&c, &m);
22     Pthread_mutex_unlock(&m);
23 }
24
25 int main(int argc, char *argv[]) {
26     printf("parent: begin\n");
27     pthread_t p;
28     Pthread_create(&p, NULL, child, NULL);
29     thr_join();
30     printf("parent: end\n");
31     return 0;
32 }
```

Figure 30.3: Parent Waiting For Child: Use A Condition Variable

# Why do we need the state variable done?

- Race condition: missed wakeup
  - Child runs immediately, signals, but no one sleeping yet
  - Parent decides to sleep and goes to sleep forever as child has already executed.

```
1 void thr_exit() {  
2     Pthread_mutex_lock(&m);  
3     Pthread_cond_signal(&c);  
4     Pthread_mutex_unlock(&m);  
5 }  
6  
7 void thr_join() {  
8     Pthread_mutex_lock(&m);  
9     Pthread_cond_wait(&c, &m);  
10    Pthread_mutex_unlock(&m);  
11 }
```

# Why do we need to lock before wait?

- Race condition: missed wakeup
  - Parent checks done to be 0, decides to sleep, interrupted
  - Child runs, sets done to 1, signals, but no one sleeping yet
  - Parent now resumes and goes to sleep forever
- Lock must be held when calling wait and signal with CV
- The wait function releases the lock before putting thread to sleep, so lock is available for signaling thread

```
1 void thr_exit() {  
2     done = 1;  
3     Pthread_cond_signal(&c);  
4 }  
5  
6 void thr_join() {  
7     if (done == 0)  
8         Pthread_cond_wait(&c);  
9 }
```