

# Operating System (CSE231)

Section-B

Quiz 4

Total Marks: 40

Time: 1 hr

---

1. (12 points) You are given the following code that inspects a file “data.txt” using system calls. Assuming that the file “data.txt” contains the content:

ABCDEFGHIJKL

Answer the following questions.

---

```
#include <stdio.h>
#include <fcntl.h>
#include <unistd.h>
#include <dirent.h>
#include <assert.h>
#include <string.h>

int main() {

    int fd1 = open("data.txt", O_RDWR);
    assert(fd1 >= 0);
    int fd2 = open("data.txt", O_RDONLY);
    assert(fd2 >= 0);
    lseek(fd1, 5, SEEK_SET);
    char buf1[6];
    char buf2[6];
    read(fd1, buf1, 5);
    buf1[5] = '\0';
    read(fd2, buf2, 5);
    buf2[5] = '\0';
    printf("fd1 read: %s\n", buf1); // A
    printf("fd2 read: %s\n", buf2); // B

    // C

    return 0;
}
```

---

- (4 points) What would be the outputs of code line //A and //B? Explain your answer.
- (4 points) Insert code at code line //C to force synchronize fd1. Also, verify whether the operation was successful.
- (4 points) In UNIX operating systems, which data structure represents a directory entry and is returned by function calls like `readdir()`? Also, what is an inode number?

Ans.

(a) fd1 read: FGHIJ

fd2 read: ABCDE

Explanation: After opening the file twice, each file descriptor receives its own open file table entry with its own independent current offset. Each open() creates a new entry with offset = 0.

fd1 seeks to offset 5, then reads 5 bytes, "FGHIJ". fd2 is still at offset 0, so it reads "ABCDE".

**Grading guide:** 1 mark each for correct output. 1 mark each for correct explanation.

(b)

---

```
int rc = write(fd1, buf1, 5);
assert(rc == 5);
rc = fsync(fd1);
assert(rc == 0);
```

---

**Grading guide:** 2 marks for mentioning fsync(fd1). 2 more marks for complete correct code.

(c) The `dirent` structure. The inode number is the index node which is an OS-level identifier.

**Grading guide:** 2 marks for each correct answer.

2. (20 points) A disk runs at 7200 RPM, and the seek time per track is 8 ms. The disk has 600 tracks, each track has 120 sectors, numbered 0–119. The **SCAN (Elevator) scheduling algorithm** is used. The disk head is initially positioned at track 320, sector 40, moving toward the higher-numbered tracks (toward track 600). The disk rotates from the lower-numbered sector to the higher-numbered sector. The following disk I/O requests arrive and must be scheduled:

Track	Sector
500	60
120	30
320	10
410	80
250	90

- (a) (3 points) Determine the order in which the requests will be serviced using the SCAN algorithm.
- (b) (15 points) For each request in that order, calculate the Seek Time, Rotational Latency/Delay, and the Total access time (this requires rotational latency per sector). Show your calculations.
- (c) (2 points) What issue might arise if a scheduler optimizes only for rotational latency and not for seek time?

Ans. Seek time = 8 ms

RPS = RPM/60 = 7200/60 = 120 revolutions/sec

Time for one rotation = 1/RPS = 1/120 = 8.3333 ms

Avg. rotational latency = revolution time/2 = 4.1667 ms

Rotational latency/sector = revolution time/sectors = 8.3333 ms/120 = 0.069444 ms/sector

(a) Given: must use Elevator SCAN moving upwards, therefore

SCAN will service requests at tracks greater than or equal to 320 in increasing track order, then reverse and service remaining lower-track requests in decreasing track order.

Requests with track greater than or equal to 320: tracks 320, 410, 500 and then reverse direction and serve lower tracks in decreasing order: remaining tracks 250 and 120.

Final service order:

(320, 10)

(410, 80)

(500, 60)

(250, 90)

(120, 30)

**Grading guide:** 3 marks for correct order. No partial marking.

(b) **Request 1:** (320, 10) starting from (320, 40)

Seek distance:  $|320 - 320| = 0$

Seek time:  $0 \times 8 \text{ ms} = 0.000 \text{ ms}$

Rotational latency: moving from sector 40 to 10. Since target < current and the rotation cannot go backwards, a full rotation is required (10+120), minus the current sector (-40). Therefore, forwardSteps =  $10 + 120 - 40 = 90$  sectors

Then, rotational latency =  $90 \times 0.069444 \text{ ms} = 6.250 \text{ ms}$

Total time = Seek + Rotational =  $0 + 6.250 = 6.250 \text{ ms}$

**Request 2:** (410, 80)

Seek distance:  $|410 - 320| = 90$  tracks

Seek time =  $90 \times 8 \text{ ms} = 720.000 \text{ ms}$

Rotational latency: sector 10 to 80

target sector > current sector, therefore: forwardSteps =  $80 - 10 = 70$  sectors

Rotational latency =  $70 \times 0.069444 \text{ ms} = 4.861 \text{ ms}$

Total time =  $720.000 + 4.861 = 724.861 \text{ ms}$

**Request 3:** (500, 60)

Seek distance:  $|500 - 410| = 90$  tracks

Seek time =  $90 \times 8 \text{ ms} = 720.000 \text{ ms}$

Rotational latency: sector 80 to 60

target < current, so forwardSteps =  $60 + 120 - 80 = 100$  sectors

Rotational latency =  $100 \times 0.069444 \text{ ms} = 6.944 \text{ ms}$

Total time =  $720.000 + 6.944 = 726.944 \text{ ms}$

**Request 4:** (250, 90) after reversing direction: track 500 to 250, sector 60 to 90

Seek distance:  $|250 - 500| = 250$  tracks

Seek time =  $250 \times 8$  ms = 2000.000 ms

Rotational latency: sector 60 to 90

forwardSteps = 90 - 60 = 30 sectors

Rotational latency =  $30 \times 0.069444$  ms = 2.083 ms

Total time =  $2000.000 + 2.083 = 2002.083$  ms

**Request 5:** (120, 30)

Seek distance:  $|120 - 250| = 130$  tracks

Seek time =  $130 \times 8$  ms = 1040.000 ms

Rotational latency: sector 90 to 30

target < current, so forwardSteps =  $30 + 120 - 90 = 60$  sectors

Rotational latency =  $60 \times 0.069444$  ms = 4.167 ms

Total time =  $1040.000 + 4.167 = 1044.167$  ms

**Grading guide:** For each request, 1 mark each for seek time, rotational latency, and total time.

(c) The scheduler might repeatedly prefer requests on the same track or nearby sectors, causing **excessive head movement (long seeks)** or **large average seek times and degraded throughput**. In the extreme, optimizing purely for rotational latency could cause **starvation** of requests far from the current head position because the scheduler always picks ones with minimal rotational wait.

**Grading guide:** 1 mark for reference to longer seek times, and 1 mark for starvation.

3. (8 points) A filesystem uses VSFS-style layout with 4 KB block size.

(a) (4 points) Each inode is 256 bytes and the inode table start address is 0 B. The disk sector size is 512 B. What is the block number and disk sector number that contain inode number 5000? Assume 1 KB = 1000 B.

(b) (4 points) Given block size 4 KB with 4-byte pointers, (i) what is the size of a file that uses 16 direct pointers and 200 pointers in a single indirect block?

(ii) What level of indirection would be required for a file of size 4.2 TB?

Ans. (a) Inode 5000's relative position =  $5000 * 256 = 1280000$  Block number =  $1280000 / 4k = 1280000 / 4000 = 320$  Sector Number =  $(1280000 + 0) / 512 = 2500$

**Grading guide:** 1 mark for correct relative position. 1 mark for correct block number. 2 marks for correct sector number.

(b) (i)  $(16 + 200) \times 4000$  B = 864000 B = 864 KB (ii) Triple indirection.

**Grading guide:** 2 marks for each part.