

File System Implementation

Revision:

- VSFS On-Disk Layout

A typical Very Simple File System (VSFS).

- Superblock: file system metadata: total inodes, total data blocks, addresses.
- Inode bitmap: 1 = allocated inode, 0 = free.
- Data bitmap: 1 = allocated block, 0 = free.
- Inode table: array of inodes; each inode ~128-256 bytes.
- Data blocks/region: user file data.

- Inodes

Each file has one inode identified by inode number.

Inode Field	Meaning
size	file size in bytes
block pointers	12 direct, 1 indirect, 1 double, 1 triple
permissions, UID, GID, timestamps	metadata

- File Size from Pointer Structure

Block size = B

Pointer size = 4 bytes; pointers per indirect block = $B / 4$

Example for 4 KB blocks, given:

Direct pointers = 12, then $12 \times 4 \text{ KB}$

Indirect block = $1 \times (1024 \times 4 \text{ KB})$

Double indirect = $1024 \times (1024 \times 4 \text{ KB})$

Triple indirect = $1024 \times (1024 \times (1024 \times 4 \text{ KB}))$

Then, total max file size:

$$= (12 + 1024 + 1024^2 + 1024^3) \times 4 \text{ KB}$$

- Directories

A directory is a file that stores: (name, inodeNumber). It includes special entries: . (dot for self) and .. (dot-dot for parent). Deletion leaves gaps and a new entry can reuse space.

- Access Paths

Steps for: *open("/foo/bar")*

- Read inode of root
- Read directory data, find "foo", get inode #
- Read inode of "foo"
- Read "foo" directory, find "bar", get inode #
- Read inode of "bar", permissions check, return file descriptor

Steps for: *read(fd)*

- Consult inode
- Find block address
- Read data block
- Update inode access time

Steps for: *write(fd)*

If write allocates a new block:

- read data bitmap, write data bitmap
- read inode, write inode
- write data block

Steps for: *File creation*

- Expensive process that involves inode bitmap, inode table, directory data, and directory inode.

- Caching & Buffering

Caching important blocks in memory:

- Static partitioning: Allocate a fixed size of memory for caching file system blocks.
- Dynamic partitioning: Allocate memory more flexibly across virtual memory and file system.

Technique	Benefit	Risk
Standard caching	Eliminates repeated I/O (static/dynamic)	cache misses

Write buffering	Batching, scheduling, avoids unnecessary writes	crash may lose recent writes
fsync()	forces durability	slower

- **Important Formulae:**

Type	Key Formulas
Inode, block / sector lookup	$\text{inodeOffset} = i \times \text{inodeSize}$ $\text{blockNo} = \text{inodeOffset} / \text{blockSize}$ $\text{sector} = (\text{inodeOffset} + \text{inodeStart}) / \text{sectorSize}$
File size from pointer model	$(\text{direct} + \text{indirect} + \text{double} + \text{triple}) \times \text{blockSize}$
Max pointers per indirect block	$\text{blockSize} / \text{pointerSize}$ (e.g., $4096 / 4 = 1024$)
Double/Triple expansion	multiply levels: 1024^2 , 1024^3

Questions:

Q1. A filesystem has block size = 4 KB, Inode size = 256 B, sector size = 512 B, and the Inode table starts at byte offset = 0. Find which block and sector contain inode number 9000.

Ans. $\text{inodeOffset} = 9000 \times 256 = 2,304,000$ bytes
 $\text{block} = 2,304,000 / 4000 = 576$
 $\text{sector} = 2,304,000 / 512 = 4500$

Q2. Block size = 4 KB. File uses 12 direct blocks + 100 pointers in its indirect block. Find file size.

Ans. $12 \times 4096 = 49,152$ bytes
 $100 \times 4096 = 409,600$ bytes
Total = 458,752 bytes

Q3. Block size = 4 KB, pointer size = 4 B. Inode supports: 12 direct + single + double + triple. What is the maximum file size this architecture can hold?

Ans. $(12 + 1024 + 1024^2 + 1024^3) \times 4$ KB

= 4,402,345,721,856 bytes \approx 4.4 TB

Q4. Given, Block size = 4096 B, pointer size = 4 B. How many pointers fit in a single indirect block?

Ans. $4096 / 4 = 1024$ pointers

Q5. For `open("/os/notes/week1.txt")`, how many inode and directory block reads (no caching) occur?

Ans.

root inode read

root dir block read -> find "os"

inode for /os read

/os directory block read -> find "notes"

inode for /os/notes read

/os/notes directory block read -> find "week1.txt"

inode for file read

Total = 7 reads

Q6. Given, Block size = 4 KB, inode size = 256 B. The inode table occupies 5 blocks. How many inodes exist?

Ans. One block holds $4096 / 256 = 16$ inodes

5 blocks: $5 \times 16 = 80$ inodes