

Lecture 9: Paging: TLB and Multi-Level Page Tables

Operating Systems

Content taken from: <https://www.cse.iitb.ac.in/~mythili/os/>

<https://pages.cs.wisc.edu/~remzi/OSTEP/>

Last few classes

- MMU does translation from virtual address to physical address
- Dynamic relocation with base + bound registers
- Segmentation
- Paging
 - Given a virtual address, how to convert to physical address?
 - Extract VPN and offset from virtual address
 - Use VPN to index into the page table (using PTBR) and get PFN
 - Concatenate PFN and offset to get the physical address

Faster Translation Using TLBs

- Add some hardware support to MMU
 - Translation-lookaside buffer (TLB)
- TLB is simply a hardware cache of popular virtual-to-physical address translations
- Upon each virtual memory reference, the hardware first checks the TLB
- If the desired translation is held in TLB, the translation is performed (quickly) without having to consult the page table (which has all translations)

TLB Contents

- A typical TLB may have 32, 64 or 128 entries
- **Fully-associative:** Hardware searches the entire TLB in parallel to find the desired translation

VPN | PFN | other bits

```
1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True) // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10     else // TLB Miss
11         PTEAddr = PTBR + (VPN * sizeof(PTE))
12         PTE = AccessMemory(PTEAddr)
13         if (PTE.Valid == False)
14             RaiseException(SEGMENTATION_FAULT)
15         else if (CanAccess(PTE.ProtectBits) == False)
16             RaiseException(PROTECTION_FAULT)
17         else
18             TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
19             RetryInstruction()
```

Figure 19.1: **TLB Control Flow Algorithm**

Example: Accessing An Array

```
int sum = 0;  
for (i = 0; i < 10; i++) {  
    sum += a[i];  
}
```

- Which of the memory accesses are going to be TLB miss and TLB hit?
 - $a[0], a[3], a[7]$ will result in a TLB miss.
 - The rest of the array elements will result in a TLB hit.




Figure 19.2: Example: An Array In A Tiny Address Space

Locality

- Spatial Locality
 - If a program accesses memory at address x , it will likely soon access memory near x .
- Temporal Locality
 - An instruction or data item that has been recently accessed will likely be re-accessed soon in the future.




Figure 19.2: Example: An Array In A Tiny Address Space

Who handles the TLB Miss?

- **Hardware-managed TLBs**
 - Hardware handles the TLB miss entirely
 - Hardware should know where the page table is stored in memory (PTBR)
 - On a miss, the hardware finds the correct PTE, extracts the desired translation, updates the TLB with the translation and retry the instruction
- **Software-managed TLBs**
 - Hardware simply raises an exception / trap
 - Pause the current instruction stream and raise the privilege level to kernel mode
 - OS Jumps to trap handler code for handling TLB misses
 - Code will find the desired translation from page table, use it to update the TLB and return from trap
 - On TLB miss, after returning from trap, the PC should point to the instruction that caused the trap and not the one after it (which is the normal execution flow).
 - There is a chance that while executing the TLB-miss handling code, an infinite chain of TLB misses can occur.
 - TLB miss handlers could be kept in physical memory (no mapping is required).
 - Reserve entries in the TLB for permanent translation slots for the handler.

TLB Issue: Context Switches

- TLB contains virtual-to-physical translations that are only valid for the currently running process
- **Obvious solution:** simply flush the TLB on context switches, thus emptying it before running the next process
- What is the problem here?

Sharing TLB across context switches

- Add an address space identifier (ASID) field in TLB
- ASID is analogous to process identifier (PID)

VPN	PFN	valid	prot	ASID
10	100	1	rwx	1
—	—	0	—	—
10	170	1	rwx	2
—	—	0	—	—

VPN	PFN	valid	prot	ASID
10	101	1	r-x	1
—	—	0	—	—
50	101	1	r-x	2
—	—	0	—	—

TLB Entry Replacement Policy

- **Option 1:** Evict the **Least Recently Used (LRU)** entry
- **Option 2:** Evict any entry **randomly**

How big can page tables be?

- Let us take an example to understand
 - 32-bit virtual address space
 - 4KB page size (2^{12})
 - 4-byte page-table entry
 - Each address location stores 1 byte of data
- For the above example,
 - How many unique addresses are possible? (2^{32})
 - What is the total size of virtual memory? ($2^{32} * 1$)
 - How many number of pages are required? ($2^{32}/2^{12} = 2^{20}$)
 - What is the size of page-table? ($2^{20} * 4 = 2^{22} = 4MB$)
 - How many bits are required for identifying a VPN? (20 bits = #pages)
 - How many bits are required for the offset within a page? ($32-20 = 12$)

Can we reduce the page table size?




Figure 20.1: A 16KB Address Space With 1KB Pages

PFN	valid	prot	present	dirty
10	1	r-x	1	0
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
23	1	rw-	1	1
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
-	0	—	—	—
28	1	rw-	1	1
4	1	rw-	1	1

Figure 20.2: A Page Table For 16KB Address Space

Multi-Level Page Tables

- Chop up the page table into page-sized units
- Don't allocate any page of page-table entries if it is entirely invalid
- Use **page directory** to track which page of page-table is valid/invalid

Multi-Level Page Tables




Figure 20.3: Linear (Left) And Multi-Level (Right) Page Tables

Multi-Level Page Tables

- **Advantages**
 - Allocates page-table space in proportion to the amount of address space used
 - Does not require the entire page table to reside contiguously in physical memory
- **Disadvantages**
 - On a TLB miss, two loads from memory will be required to get the PFN information for data
 - One for the page directory (PDE) and another for the PTE
 - More complex page-table lookup
- **Time-Space Tradeoff**
- Compare the above advantages and disadvantages with respect to linear page tables

Multi-Level Page Table: Example

- 16KB virtual address space
- 64-byte pages
- PTE is 4 bytes in size
- How many number of pages will be there in total? ($16\text{KB}/64 = 2^{14}/2^6 = 2^8$)

How many bits are required for fully addressing this address space? (14 bits)

- How many bits for VPN? (8 bits)
- How many bits for offset? (6 bits)

0000 0000	code
0000 0001	code
0000 0010	(free)
0000 0011	(free)
0000 0100	heap
0000 0101	heap
0000 0110	(free)
0000 0111	(free)
.....	... all free ...
1111 1100	(free)
1111 1101	(free)
1111 1110	stack
1111 1111	stack

Figure 20.4: A 16KB Address Space With 64-byte Pages

Multi-Level Page Table: Example


- What will be the number of entries in a linear page table? ($2^8 = \# \text{pages}$)
- What is the size of linear page table? ($2^8 * 4 = 2^{10} \text{ bytes}$)
- How many pages will be there in the linear page table? ($2^{10} / 64 = 2^4$)
- How many PTEs can be stored in one page? ($64 / 4 = 16$)
- How many entries will be there in the page directory? ($2^4 = \# \text{pages in the page table}$)

0000 0000	code
0000 0001	code
0000 0010	(free)
0000 0011	(free)
0000 0100	heap
0000 0101	heap
0000 0110	(free)
0000 0111	(free)
.....	... all free ...
1111 1100	(free)
1111 1101	(free)
1111 1110	stack
1111 1111	stack

Figure 20.4: A 16KB Address Space With 64-byte Pages

Multi-Level Page Table: Example


- What is the number of bits required to uniquely identify an entry in the page directory? Or what is the number of bits required for indexing into the page directory?



- $\text{PDEAddr} = \text{PageDirBase} + (\text{PDIIndex} * \text{sizeof(PDE)})$

Multi-Level Page Table: Example

- Once we have the PDEAddr, we can get the address of the required page of the page table.
- What is the number of bits required to uniquely identify an entry in a page of the page table? Or what is the number of bits required for indexing into the page of the page table?



- $\text{PTEAddr} = (\text{PDE.PFN} \ll \text{SHIFT}) + (\text{PTIndex} * \text{sizeof(PTE)})$

Multi-Level Page Table: Example

- Translate the virtual address **11 1111 1000 0000** to physical address.
 - Top 4 bits (1111) for PDIIndex. Lookup entry 15. This is 101.
 - Next 4 bits (1110) is for PTIndex. Lookup entry 14. This is 55.

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
101	1	—	0	—	55	1	rw-
					45	1	rw-

Figure 20.5: A Page Directory, And Pieces Of Page Table

```

1  VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2  (Success, TlbEntry) = TLB_Lookup(VPN)
3  if (Success == True) // TLB Hit
4      if (CanAccess(TlbEntry.ProtectBits) == True)
5          Offset = VirtualAddress & OFFSET_MASK
6          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7          Register = AccessMemory(PhysAddr)
8      else
9          RaiseException(PROTECTION_FAULT)
10     else // TLB Miss
11         // first, get page directory entry
12         PDIndex = (VPN & PD_MASK) >> PD_SHIFT
13         PDEAddr = PDBR + (PDIndex * sizeof(PDE))
14         PDE = AccessMemory(PDEAddr)
15         if (PDE.Valid == False)
16             RaiseException(SEGMENTATION_FAULT)
17         else
18             // PDE is valid: now fetch PTE from page table
19             PTIndex = (VPN & PT_MASK) >> PT_SHIFT
20             PTEAddr = (PDE.PFN << SHIFT) + (PTIndex * sizeof(PTE))
21             PTE = AccessMemory(PTEAddr)
22             if (PTE.Valid == False)
23                 RaiseException(SEGMENTATION_FAULT)
24             else if (CanAccess(PTE.ProtectBits) == False)
25                 RaiseException(PROTECTION_FAULT)
26             else
27                 TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
28             RetryInstruction()

```

Figure 20.6: Multi-level Page Table Control Flow

Inverted Page Tables

- A single page table that has an entry for each physical page of the system
- Each entry tells
 - Which process is using this physical page?
 - Which virtual page of that process maps to this physical page?
- Given a VPN, how do we find the PFN?
 - Linear scan
 - Hash tables

Swapping the Page Tables to Disk

- Swap some of the page tables to disk when low on physical memory