

Lecture 13: Concurrency: Threads and Locks

Operating Systems

Content taken from: <https://pages.cs.wisc.edu/~remzi/OSTEP/>

Till now

- CPU virtualization
 - Mechanism
 - Policy
- Memory virtualization
 - Mechanism
 - Segmentation
 - Paging
 - Swapping
 - Policy
 - LRU
 - Random

Concurrency

- Till now, we assumed that each process has a single point of execution
 - A single PC where instructions are being fetched from and executed
- Now, we introduce a new abstraction: **thread**
- Each process can have multiple threads i.e. multiple points of execution
 - Multiple PCs
 - Such a process is called **multi-threaded** process

What is a thread?

- Each thread is nothing but a separate process
- All the threads belonging to the same process share the same address space
 - Code, Data, Heap is shared but not the stack
- Each thread has a its own PC, private set of registers and stack
- Context switch between threads is similar to context switch between processes
 - OS needs to save thread's state to **thread control block (TCB)**
 - OS does not need to switch the page table. Why?





Figure 26.1: Single-Threaded And Multi-Threaded Address Spaces

Why use Threads?

- **Parallelism**
 - If you are running a program on a multi-processor system, you may want to utilize all the processors simultaneously
 - E.g. Adding two large arrays
- **Avoiding blocking**
 - Threading enables overlap of I/O with other activities ***within*** a single program
 - E.g. While a thread is doing I/O, another thread can do some computations or processing

Example: Thread Creation

```
1 #include <stdio.h>
2 #include <assert.h>
3 #include <pthread.h>
4 #include "common.h"
5 #include "common_threads.h"
6
7 void *mythread(void *arg) {
8     printf("%s\n", (char *) arg);
9     return NULL;
10 }
11
12 int
13 main(int argc, char *argv[]) {
14     pthread_t p1, p2;
15     int rc;
16     printf("main: begin\n");
17     Pthread_create(&p1, NULL, mythread, "A");
18     Pthread_create(&p2, NULL, mythread, "B");
19     // join waits for the threads to finish
20     Pthread_join(p1, NULL);
21     Pthread_join(p2, NULL);
22     printf("main: end\n");
23     return 0;
24 }
```

Figure 26.2: Simple Thread Creation Code (`t0.c`)

main	Thread 1	Thread2
starts running		
prints "main: begin"		
creates Thread 1	runs	
creates Thread 2	prints "A"	
waits for T1	returns	
waits for T2		
	runs	
	prints "B"	
	returns	
prints "main: end"		

Figure 26.3: Thread Trace (1)

main	Thread 1	Thread2
starts running		
prints “main: begin”		
creates Thread 1		
	runs	
	prints “A”	
	returns	
creates Thread 2		
		runs
		prints “B”
		returns
waits for T1		
<i>returns immediately; T1 is done</i>		
waits for T2		
<i>returns immediately; T2 is done</i>		
prints “main: end”		

Figure 26.4: Thread Trace (2)

main	Thread 1	Thread2
starts running		
prints “main: begin”		
creates Thread 1		runs
creates Thread 2		prints “B”
		returns
waits for T1		
	runs	
	prints “A”	
	returns	
waits for T2		
<i>returns immediately; T2 is done</i>		
prints “main: end”		

Figure 26.5: Thread Trace (3)

Example: Shared Data

```
1 #include <stdio.h>
2 #include <pthread.h>
3 #include "common.h"
4 #include "common_threads.h"
5
6 static volatile int counter = 0;
7
8 // mythread()
9 //
10 // Simply adds 1 to counter repeatedly, in a loop
11 // No, this is not how you would add 10,000,000 to
12 // a counter, but it shows the problem nicely.
13 //
14 void *mythread(void *arg) {
15     printf("%s: begin\n", (char *) arg);
16     int i;
17     for (i = 0; i < 1e7; i++) {
18         counter = counter + 1;
19     }
20     printf("%s: done\n", (char *) arg);
21     return NULL;
22 }
```

```
24 // main()
25 //
26 // Just launches two threads (pthread_create)
27 // and then waits for them (pthread_join)
28 //
29 int main(int argc, char *argv[]) {
30     pthread_t p1, p2;
31     printf("main: begin (counter = %d)\n", counter);
32     Pthread_create(&p1, NULL, mythread, "A");
33     Pthread_create(&p2, NULL, mythread, "B");
34
35     // join waits for the threads to finish
36     Pthread_join(p1, NULL);
37     Pthread_join(p2, NULL);
38     printf("main: done with both (counter = %d)\n",
39             counter);
40
41 }
```

Figure 26.6: Sharing Data: Uh Oh (**t1.c**)

Indeterminate Result

```
prompt> gcc -o main main.c -Wall -pthread; ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 20000000)
```

```
prompt> ./main
main: begin (counter = 0)
A: begin
B: begin
A: done
B: done
main: done with both (counter = 19345221)
```

Uncontrolled Scheduling

- In the previous example, assembly code which updates the counter looks like the following:

mov 0x8049a1c, %eax	100	mov 0x8049a1c, %eax
add \$0x1, %eax	105	add \$0x1, %eax
mov %eax, 0x8049a1c	108	mov %eax, 0x8049a1c

- The above piece of code is called as **critical section**
- A critical section is a piece of code that accesses a shared variable and must not be concurrently executed by more than one thread.

Race Condition

OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
	<i>before critical section</i>		100	0	50
	mov 8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
	<i>save T1</i>				
	<i>restore T2</i>		100	0	50
		mov 8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 8049a1c	113	51	51
interrupt					
	<i>save T2</i>				
	<i>restore T1</i>		108	51	51
	mov %eax, 8049a1c		113	51	51

Figure 26.7: The Problem: Up Close and Personal

How to avoid race conditions?

- We want **mutual exclusion**
- If one thread is executing within the critical section, the others should be prevented from doing so.
- In other words, we should write multi-threaded code that accesses critical sections in a **synchronized** and **controlled** manner
 - This avoids race conditions
 - Results in deterministic program outputs

Locks and Pthread Locks

- Programmers annotate source code with locks, putting them around critical sections, and thus ensure that any such critical section executes as if it were a **single atomic instruction**

```
1 lock_t mutex; // some globally-allocated lock 'mutex'  
2 ...  
3 lock(&mutex);  
4 balance = balance + 1;  
5 unlock(&mutex);  
  
1 pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2  
3 Pthread_mutex_lock(&lock); // wrapper; exits on failure  
4 balance = balance + 1;  
5 Pthread_mutex_unlock(&lock);
```

Evaluating Locks

- **Correctness:** Does the lock ensure mutual exclusion?
- **Fairness:** Does each thread contending for the lock get a fair shot at acquiring it once it is free?
- **Performance:** How much time overhead does the lock add?

How can we build a lock?

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 -> lock is available, 1 -> held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11      mutex->flag = 1;           // now SET it!
12  }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

Figure 28.1: First Attempt: A Simple Flag

Lock using a simple flag does not work

- Mutual Exclusion is not ensured

Thread 1	Thread 2
call lock ()	
while (flag == 1)	
interrupt: switch to Thread 2	
	call lock ()
	while (flag == 1)
	flag = 1;
	interrupt: switch to Thread 1
flag = 1; // set flag to 1 (too!)	

Figure 28.2: Trace: No Mutual Exclusion

- Performance overhead: **Spin-waiting**

Building Working Spin Locks with Test-And-Set

```
1 int TestAndSet(int *old_ptr, int new) {
2     int old = *old_ptr; // fetch old value at old_ptr
3     *old_ptr = new;    // store 'new' into old_ptr
4     return old;        // return the old value
5 }
```

```
1     typedef struct __lock_t {
2         int flag;
3     } lock_t;
4
5     void init(lock_t *lock) {
6         // 0: lock is available, 1: lock is held
7         lock->flag = 0;
8     }
9
10    void lock(lock_t *lock) {
11        while (TestAndSet(&lock->flag, 1) == 1)
12            ; // spin-wait (do nothing)
13    }
14
15    void unlock(lock_t *lock) {
16        lock->flag = 0;
17    }
```

Figure 28.3: A Simple Spin Lock Using Test-and-set

Evaluating Spin Locks

- Correctness?
- Fairness?
- Performance?