

Lecture 19: Files and Directories

Operating Systems

Content taken from: <https://pages.cs.wisc.edu/~remzi/OSTEP/>

<https://www.cse.iitb.ac.in/~mythili/os/>

File and Directory as Abstractions

- Two key abstractions for virtualizing persistent storage
- **File** is a linear array of bytes, stored persistently
 - File name (human readable)
 - OS-level identifier (**inode number**)
- **Directory** contains other subdirectories and files, along with their inode numbers
 - Stored like a file, whose contents are filename-to-inode mappings

Directory Tree

- Root directory
- **Separator** for naming subsequent sub-directories
- **Absolute pathname**
- **File Type**




Figure 39.1: An Example Directory Tree

Creating a File

- **open()** system call

```
int fd = open ("foo", O_CREAT|O_WRONLY|O_TRUNC,  
               S_IRUSR|S_IWUSR);
```

- Returns a **file descriptor**
 - File descriptor is just an integer, private per process
 - Used to access a file and perform operations on it
 - Think of file descriptor as a pointer to an object of type file
 - File descriptor for
 - STD_IN is 0
 - STD_OUT is 1
 - STD_ERR is 2
- What happens when we create or open a file?
 - An entry is created in the system-wide **Open File Table (OFT)**.
 - Each process keeps an array of file descriptors each of which refers to an entry in the OFT.

Reading and Writing Files

- `Read(file_descriptor, buffer_to_read_into, buffer_size)`
 - Returns the number of bytes successfully read
- `Write(file_descriptor, buffer_to_write_from, bytes_to_write)`
 - Returns the number of bytes successfully written
- Reading and writing happens **sequentially** by default
 - Successive read/write calls fetch from current offset

Examples

System Calls	Return Code	Current Offset
fd = open("file", O_RDONLY);	3	0
read(fd, buffer, 100);	100	100
read(fd, buffer, 100);	100	200
read(fd, buffer, 100);	100	300
read(fd, buffer, 100);	0	300
close(fd);	0	—

System Calls	Return Code	OFT[10] Current Offset	OFT[11] Current Offset
fd1 = open("file", O_RDONLY);	3	0	—
fd2 = open("file", O_RDONLY);	4	0	0
read(fd1, buffer1, 100);	100	100	0
read(fd2, buffer2, 100);	100	100	100
close(fd1);	0	—	100
close(fd2);	0	—	—

Reading And Writing, But Not Sequentially

- Sometimes, however, it is useful to be able to read or write to a specific offset within a file
- `Iseek()` system call be used to perform read or write operations at **random** offsets within the document

```
off_t lseek(int fildes, off_t offset, int whence);
```

If whence is `SEEK_SET`, the offset is set to offset bytes.

If whence is `SEEK_CUR`, the offset is set to its current location plus offset bytes.

If whence is `SEEK_END`, the offset is set to the size of the file plus offset bytes.

Example

System Calls	Return Code	Current Offset
fd = open("file", O_RDONLY);	3	0
lseek(fd, 200, SEEK_SET);	200	200
read(fd, buffer, 50);	50	250
close(fd);	0	-

Shared File Table Entries

- Two processes accessing the same file at the same time
 - Each will have its own entry in the OFT
 - Each process can independently perform read/write on the file.
- In some cases, an entry in OFT is **shared**.
- **Dup()** system call

```
int main(int argc, char *argv[]) {
    int fd = open("README", O_RDONLY);
    assert(fd >= 0);
    int fd2 = dup(fd);
    // now fd and fd2 can be used interchangeably
    return 0;
}
```

Figure 39.4: Shared File Table Entry With **dup ()** (**dup.c**)

Parent and child process sharing file table entry

```
int main(int argc, char *argv[]) {
    int fd = open("file.txt", O_RDONLY);
    assert(fd >= 0);
    int rc = fork();
    if (rc == 0) {
        rc = lseek(fd, 10, SEEK_SET);
        printf("child: offset %d\n", rc);
    } else if (rc > 0) {
        (void) wait(NULL);
        printf("parent: offset %d\n",
               (int) lseek(fd, 0, SEEK_CUR));
    }
    return 0;
}
```

Figure 39.2: Shared Parent/Child File Table Entries (`fork-seek.c`)

```
prompt> ./fork-seek
child: offset 10
parent: offset 10
prompt>
```




Figure 39.3: Processes Sharing An Open File Table Entry

Writing Immediately with fsync()

- Writes performed using **write()** are buffered in memory for some time for performance reasons.
- The writes are issued to the storage device at a later point in time.
- We can force immediate write to storage device using **fsync()**

```
int fd = open("foo", O_CREAT|O_WRONLY|O_TRUNC,  
              S_IRUSR|S_IWUSR);  
assert(fd > -1);  
int rc = write(fd, buffer, size);  
assert(rc == size);  
rc = fsync(fd);  
assert(rc == 0);
```

Other system calls

- **rename(“new_file_name”, “old_file_name”)** for renaming a file
- **fstat(file_descriptor)** for fetching meta data associated with a file
 - Metadata is stored in a structure called as **inode**.
 - **Inode** is a persistent data structure stored on the disk.

```
struct stat {  
    dev_t      st_dev;      // ID of device containing file  
    ino_t      st_ino;      // inode number  
    mode_t     st_mode;     // protection  
    nlink_t    st_nlink;    // number of hard links  
    uid_t      st_uid;      // user ID of owner  
    gid_t      st_gid;      // group ID of owner  
    dev_t      st_rdev;     // device ID (if special file)  
    off_t      st_size;     // total size, in bytes  
    blksize_t  st_blksize;  // blocksize for filesystem I/O  
    blkcnt_t   st_blocks;   // number of blocks allocated  
    time_t     st_atime;    // time of last access  
    time_t     st_mtime;    // time of last modification  
    time_t     st_ctime;    // time of last status change  
};
```

Figure 39.5: The **stat** structure.

Directory-related System calls

- **mkdir("file_name", permissions)** for creating a directory
- **rmdir()** for deleting a empty directory
- **Reading Directory**

```
int main(int argc, char *argv[]) {  
    DIR *dp = opendir(".");  
    assert(dp != NULL);  
    struct dirent *d;  
    while ((d = readdir(dp)) != NULL) {  
        printf("%lu %s\n", (unsigned long) d->d_ino,  
               d->d_name);  
    }  
    closedir(dp);  
    return 0;  
}  
  
struct dirent {  
    char          d_name[256]; // filename  
    ino_t         d_ino;      // inode number  
    off_t         d_off;      // offset to the next dirent  
    unsigned short d_reclen; // length of this record  
    unsigned char  d_type;   // type of file  
};
```