

# Tutorial 3

---

Process Scheduling



INDRAPRASTHA INSTITUTE *of*  
INFORMATION TECHNOLOGY **DELHI**



# What is Process Scheduling?

---

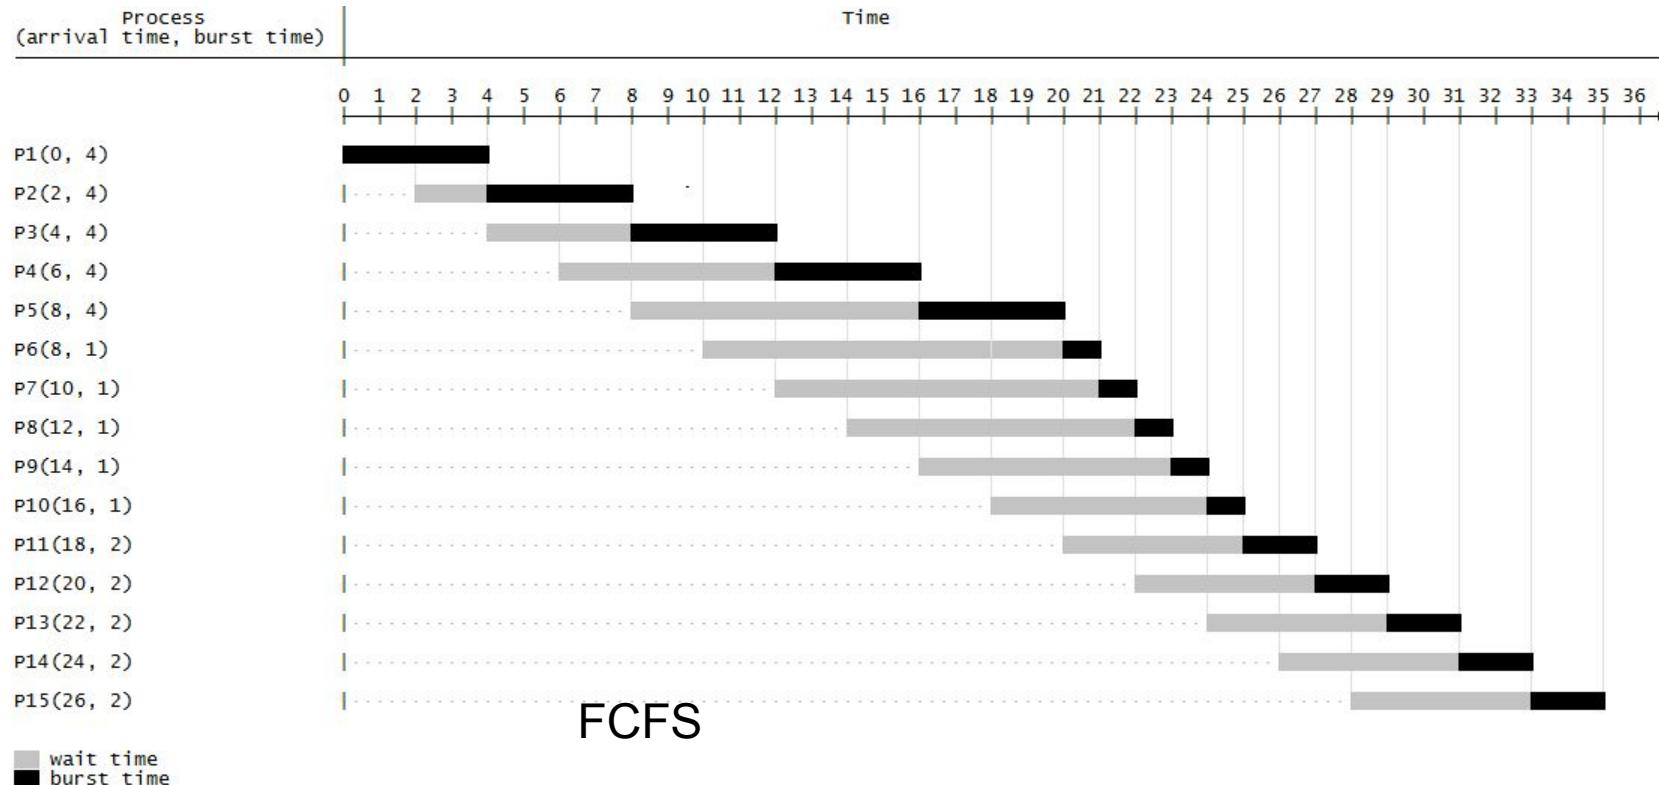
- Modern Operating Systems run more than one process ‘at a time’ (more like thousands of processes at a time),
- What do you do if the number of simultaneously running processes on a system exceed the number of physical CPU cores?
- You need a time sharing system which “switches” between processes very rapidly to give the illusion of simultaneity.
- Process Scheduling is all about how this “switching” between processes happens in the OS.

# Scheduling Policies?

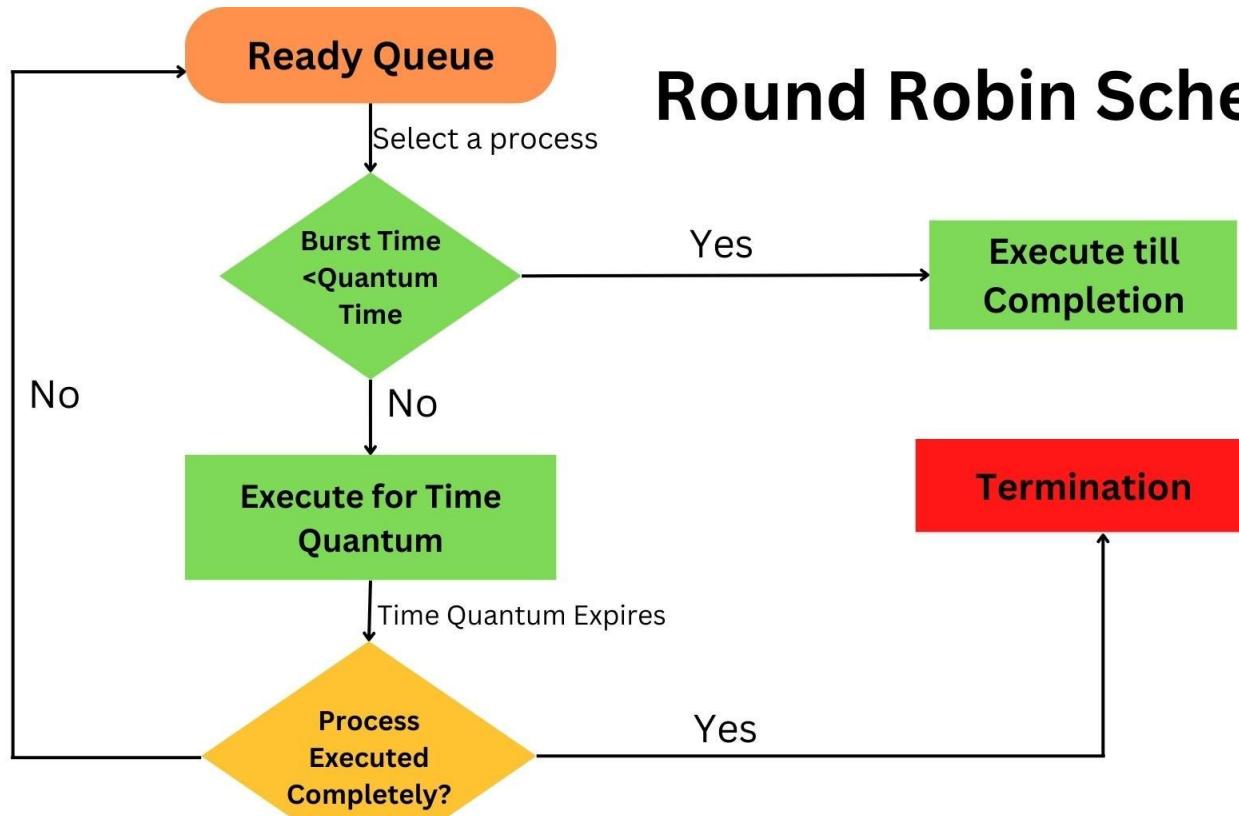
---

- First Come First Serve
- Round Robin
- $O(n)$
- $O(1)$
- CFS (Current default Linux Scheduler)

# Scheduling Policies

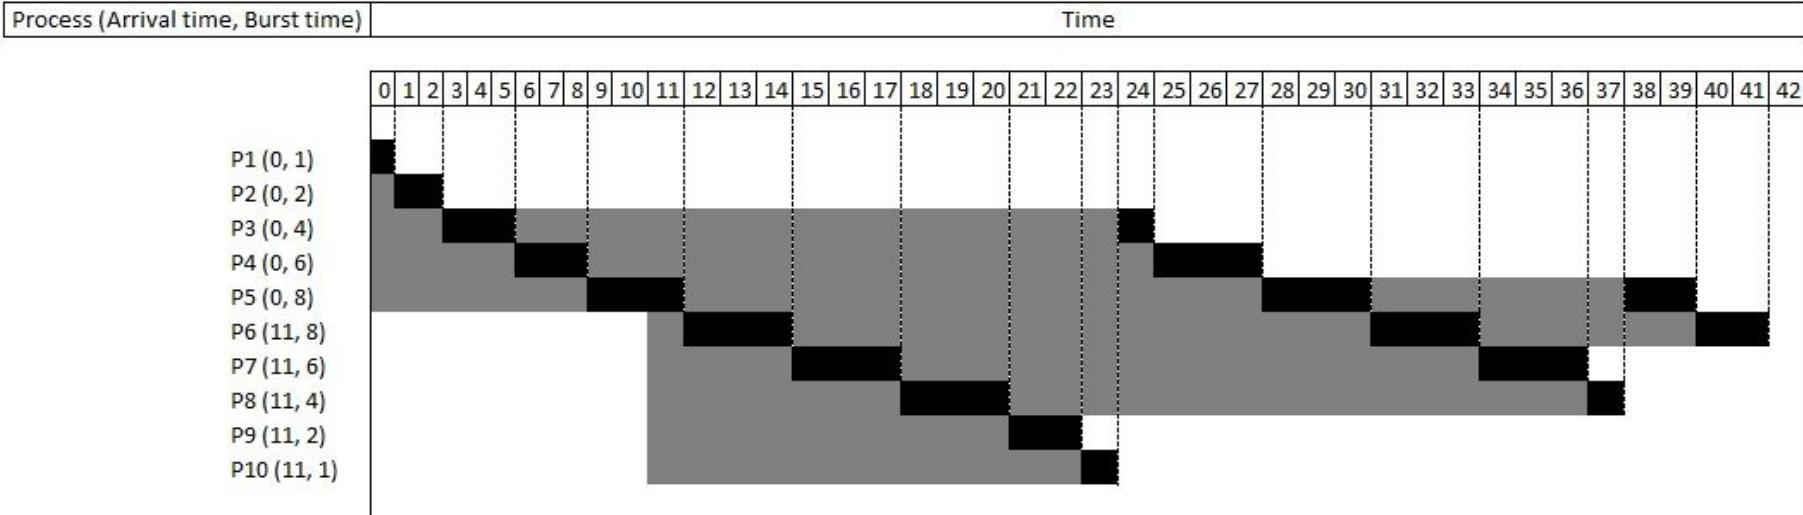


# Scheduling Policies



## Round Robin Scheduling

# Scheduling Policies

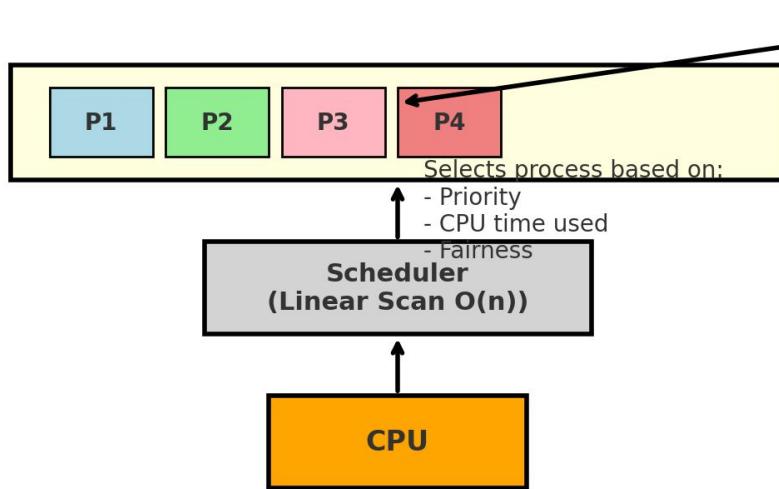


Quantum = 3

Wait time  
Burst time

Round Robin

# Scheduling Policies



**O(n) Scheduler**

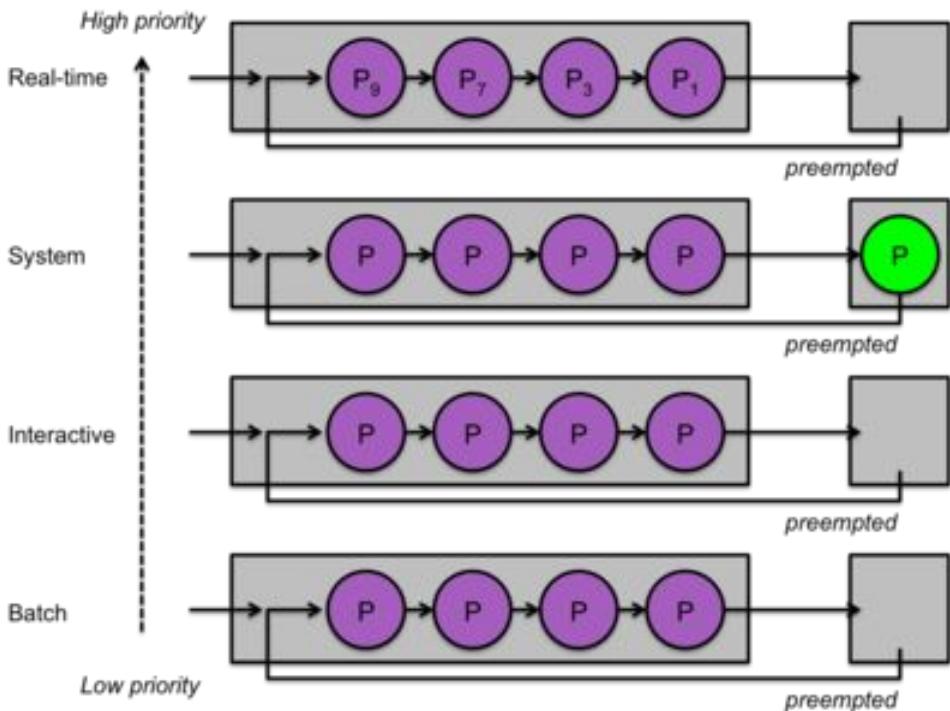
At every context switch

- Scan the list of runnable processes
- Compute priorities
- Select the best process to run

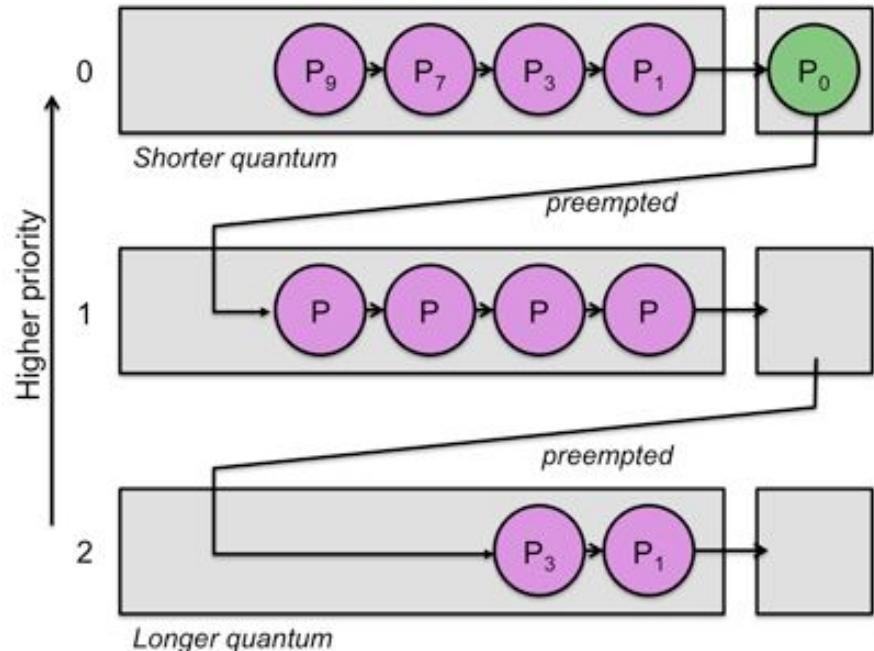
$O(n)$ , when  $n$  is the number of runnable processes ... Not scalable!!

Used in Linux Kernel 2.4 - 2.6.

# Scheduling Policies



Multi-Level Queue



Multi-Level Feedback Queue

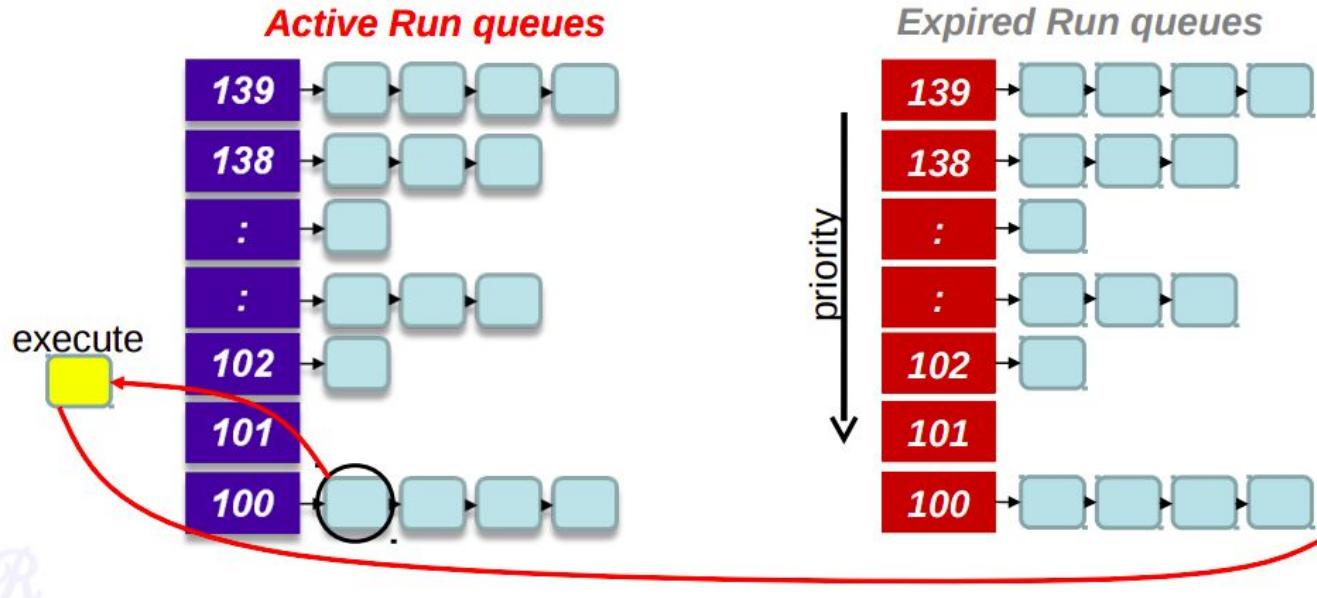
# Scheduling Policies

Base Priority = 120



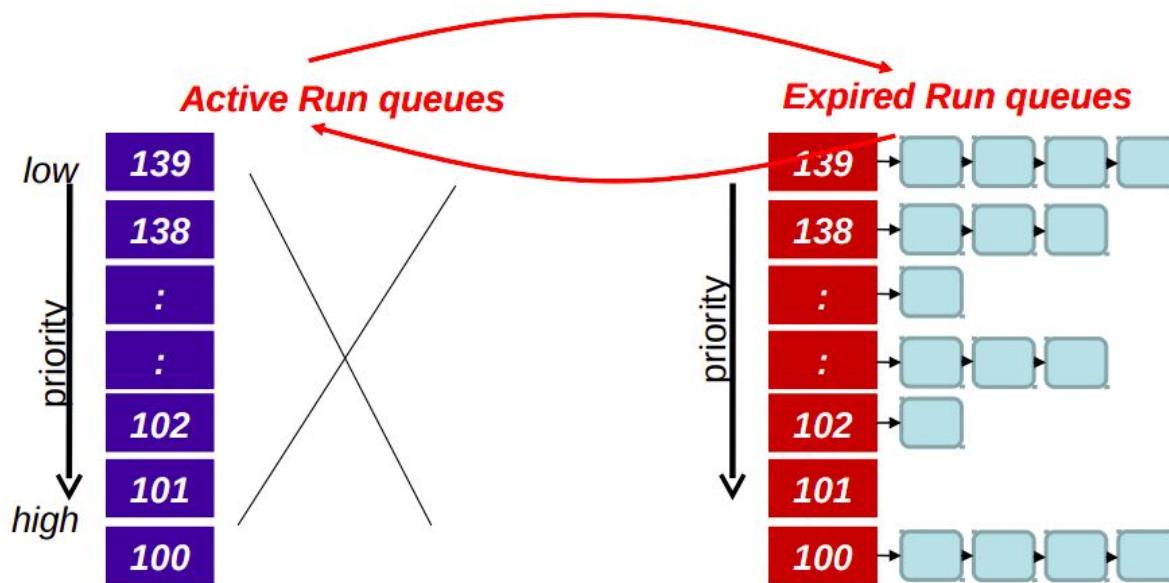
O(1) Scheduler

# Scheduling Policies



$O(1)$  Scheduler

# Scheduling Policies



O(1) Scheduler

Used in Linux 2.6-2.6.22

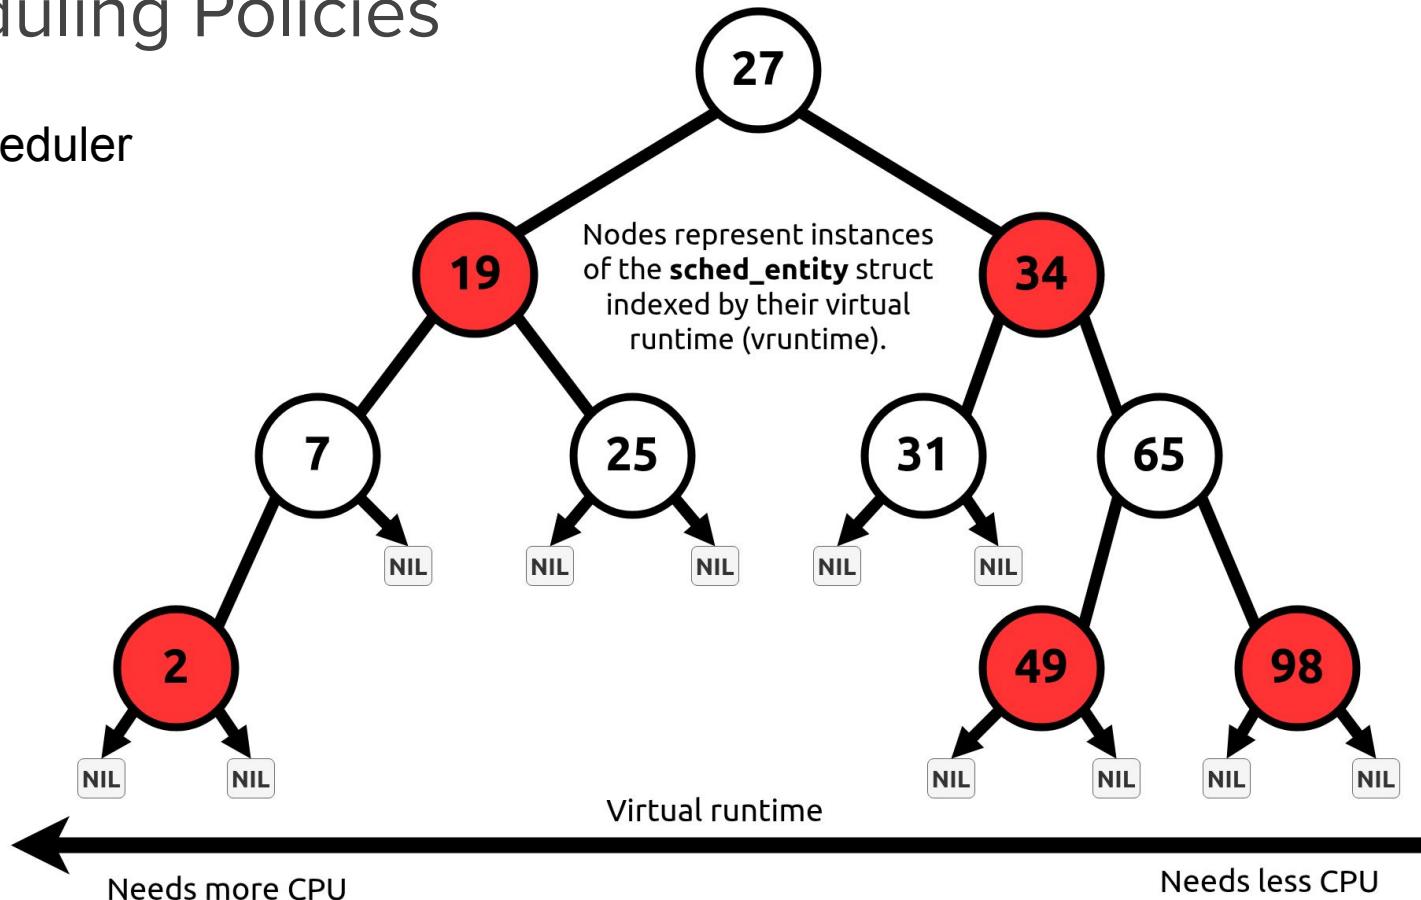
# Scheduling Policies - CFS

---

- Target Scheduler Latency ( $tl = 6 - 24$  ms)
  - If there are  $n$  runnable processes, then each process will execute for  $(tl / n)$  ms.
- With each runnable process is included a virtual runtime (**vruntime**)
  - At every scheduling point, if process has run for  $t$  ms, then **vruntime** +=  $t$
- When timer interrupt occurs
  - Choose the task with the lowest vruntime (**min\_vruntime**)
  - Compute its dynamic timeslice ( $tl/n$ ) and run
- When interrupt occurs again, context switch with another task with a smaller runtime

# Scheduling Policies

## CFS Scheduler



# Process Priority

---

- When multiple processes simultaneously require CPU time, the system's scheduling policy and process CPU priorities determine which processes get it.
- High priority processes are scheduled before lower priority processes on the CPU and are given more CPU time.

# ‘Niceness’ of a process

---

- At the user level, the priority of processes is defined in terms of their nice value.
- Nice values range from -20 (highest priority) to 19 (lowest priority).
- The nice value can be interpreted as how ‘nice’ a process is towards other processes in terms of giving up CPU time.

# Command Line Tools

# htop

# htop

---

- htop command in Linux system is a command line utility that allows the user to interactively monitor the system's vital resources or server's processes in real time.
- We can observe all processes running on the system, along with their command line arguments, select multiple processes and act on them all at once.
- htop also prints full command lines for processes and allows one to scroll both vertically and horizontally for processes and command lines respectively.

[htop command in Linux with examples - GeeksforGeeks](#)

nice/renice

# nice/renice

---

- nice command helps in execution of a process with modified scheduling priority. If we give a process a higher priority, then Kernel will allocate more CPU time to that process..
- renice command allows you to change and modify the scheduling priority of an already running process. Linux Kernel schedules the process and allocates CPU time accordingly for each of them.
- **nice [OPTION] [COMMAND [ARG]...]**
- **renice [-n] priority [-gl-pl-u] identifier...**

[Nice and Renice Command in Linux with Examples - GeeksforGeeks](#)

chrt

# chrt

---

- chrt command in Linux is known for manipulating the real-time attributes of a process.
- It sets or retrieves the real-time scheduling attributes of an existing PID, or runs the command with the given attributes.
- **chrt [options] priority command argument ...**
- **chrt [options] -p [priority] PID**

[chrt command in Linux with examples - GeeksforGeeks](#)

# Scheduling Policy Options

---

- SCHED\_BATCH : Use Scheduling batch processes algorithm.
- SCHED\_FIFO : Uses First In-First Out scheduling algorithm. This scheduling method is used on Batch-Systems, it is NON-PREEMPTIVE. It implements just one queue which holds the tasks in the order they come in.
- SCHED\_IDLE: Used for running very low priority background jobs.
- SCHED\_OTHER: Uses Default Linux time-sharing scheduling algorithm or simply the standard round-robin time-sharing policy.
- SCHED\_RR Uses Round Robin scheduling algorithm and is used as the default algorithm if not specified. It is an algorithm used for PREEMPTIVE scheduling.

[chrt command in Linux with examples - GeeksforGeeks](#)

# C Library Functions

# C Library Functions

---

- int sched\_setscheduler (pid\_t pid, int policy, const struct sched\_param \*param)
- int sched\_getscheduler (pid\_t pid)
- int sched\_setparam (pid\_t pid, const struct sched\_param \*param)
- int sched\_getparam (pid\_t pid, struct sched\_param \*param)

[https://www.gnu.org/software/libc/manual/html\\_node/Basic-Scheduling-Functions.html](https://www.gnu.org/software/libc/manual/html_node/Basic-Scheduling-Functions.html)

# Questions